CO Open in Colab

Unsupervised Sentiment Analysis with Lexicon Models

We talked about unsupervised learning methods in the past, which refer to specific modeling methods that can be applied directly to data features without the presence of labeled data. One of the major challenges in any organization is getting labeled datasets due the lack of time as well as resources to do this tedious task. Unsupervised methods are very useful in this scenario and we look at some of these methods in this section. Even though we have labeled data, this section should give you a good idea of how lexicon based models work and you can apply them to your own datasets when you do not have labeled data.

Unsupervised sentiment analysis models use well curated knowledgebases, ontologies, lexicons, and databases, which have detailed information pertaining to subjective words, phrases including sentiment, mood, polarity, objectivity, subjectivity, and so on.

A lexicon model typically uses a lexicon, also known as a dictionary or vocabulary of words specifically aligned to sentiment analysis. These lexicons contain a list of words associated with positive and negative sentiment, polarity (magnitude of negative or positive score), parts of speech (POS) tags, subjectivity classifiers (strong, weak, neutral), mood, modality, and so on.

You can use these lexicons and compute the sentiment of a text document by matching the presence of specific words from the lexicon and then looking at other factors like presence of negation parameters, surrounding words, overall context, phrases, and aggregate overall sentiment polarity scores to decide the final sentiment score.

There are several popular lexicon models used for sentiment analysis. Some of them are as follows:

- Bing Liu's lexicon
- MPQA subjectivity lexicon
- Pattern lexicon
- TextBlob lexicon
- AFINN lexicon
- SentiWordNet lexicon
- VADER lexicon

This is not an exhaustive list of lexicon models but these are definitely among the most popular ones available today.

▾ Install Dependencies

```
!pip install textblob
!pip install textsearch
!pip install contractions
```

```
!pip install afinn
import nltk
nltk.download('punkt')
nltk.download('stopwords')
nltk.download('vader_lexicon')
```

```
    Requirement already satisfied: textblob in /usr/local/lib/python3.7/dist-packages (0.15.3)
    Requirement already satisfied: nltk>=3.1 in /usr/local/lib/python3.7/dist-packages (from textblob) (3.2.5)
    Requirement already satisfied: six in /usr/local/lib/python3.7/dist-packages (from nltk>=3.1->textblob) (1.15.0)
    Requirement already satisfied: textsearch in /usr/local/lib/python3.7/dist-packages (0.0.21)
    Requirement already satisfied: pyahocorasick in /usr/local/lib/python3.7/dist-packages (from textsearch) (1.4.4)
    Requirement already satisfied: anyascii in /usr/local/lib/python3.7/dist-packages (from textsearch) (0.3.0)
    Requirement already satisfied: contractions in /usr/local/lib/python3.7/dist-packages (0.1.66)
    Requirement already satisfied: textsearch>=0.0.21 in /usr/local/lib/python3.7/dist-packages (from contractions) (0.0.21)
    Requirement already satisfied: pyahocorasick in /usr/local/lib/python3.7/dist-packages (from textsearch>=0.0.21->contractions) (1.4.4)
    Requirement already satisfied: anyascii in /usr/local/lib/python3.7/dist-packages (from textsearch>=0.0.21->contractions) (0.3.0)
    Requirement already satisfied: afinn in /usr/local/lib/python3.7/dist-packages (0.1)
    [nltk_data] Downloading package punkt to /root/nltk_data...
    [nltk_data]   Package punkt is already up-to-date!
    [nltk_data] Downloading package stopwords to /root/nltk_data...
    [nltk_data]   Package stopwords is already up-to-date!
    [nltk_data] Downloading package vader_lexicon to /root/nltk_data...
    [nltk_data]   Package vader_lexicon is already up-to-date!
    True
```

▾ Load Dependencies

```
import pandas as pd
import numpy as np
import nltk
import textblob
from sklearn.metrics import confusion_matrix, classification_report
np.set_printoptions(precision=2, linewidth=80)
```

▾ Load Dataset

```
# adapted code to download files from storage
from google.colab import files

uploaded = files.upload()
```

<div style="border:1px solid #ccc; padding:4px;">Choose Files</div> noneutral_o…eb2022.xlsx

- **noneutral_orange_solana_15feb2022.xlsx**(application/vnd.openxmlformats-officedocument.spreadsheetml.sheet) - 51971 bytes, last modified: 2/18/2022 - 100% done
  Saving noneutral_orange_solana_15feb2022.xlsx to noneutral_orange_solana_15feb2022.xlsx

```
# adapted code to read xlsx files
import pandas as pd
import io
dataset = pd.read_excel(io.BytesIO(uploaded['noneutral_orange_solana_15feb2022.xlsx']))
```

```
#reviews = np.array(dataset['review'])
reviews = np.array(dataset['text']) # the dataset field is called 'text' on orange xlsx results
sentiments = np.array(dataset['sentiment'])

# extract data for model evaluation
test_reviews = reviews[1:]
test_sentiments = sentiments[1:]
sample_review_ids = [2,5,10,50,100]
```

▾ Sentiment Analysis with TextBlob

The lexicon that TextBlob uses is the same one as pattern and is available in their source code on GitHub (https://github.com/sloria/TextBlob/blob/dev/textblob/en/en-sentiment.xml). Some sample examples are shown from the lexicon as follows.

```
<word form="abhorrent" wordnet_id="a-1625063" pos="JJ" sense="offensive
to the mind" polarity="-0.7" subjectivity="0.8" intensity="1.0"
reliability="0.9" />
<word form="able" cornetto_synset_id="n_a-534450" wordnet_id="a-01017439"
pos="JJ" sense="having a strong healthy body" polarity="0.5"
subjectivity="1.0" intensity="1.0" confidence="0.9" />
```

Typically, specific adjectives have a polarity score (negative/positive, -1.0 to +1.0) and a subjectivity score (objective/subjective, +0.0 to +1.0).

The reliability score specifies if an adjective was hand-tagged (1.0) or inferred (0.7). Words are tagged per sense, e.g., ridiculous (pitiful) = negative, ridiculous (humorous) = positive.

The Cornetto id (lexical unit id) and Cornetto synset id refer to the Cornetto lexical database for Dutch. The WordNet id refers to the WordNet3 lexical database for English. The part-of-speech tags (POS) use the Penn Treebank convention. Let's look at how we can use TextBlob for sentiment analysis.

Predict sentiment for sample reviews

```
for review, sentiment in zip(test_reviews[sample_review_ids], test_sentiments[sample_review_ids]):
    print('REVIEW:', review)
    print('Actual Sentiment:', sentiment)
    print('Predicted Sentiment polarity:', textblob.TextBlob(review).sentiment.polarity)
    print('-'*60)
```

```
REVIEW: @Brewtoshi @solana @brave @anniekatrina @BrendanEich ^^^^ scroll on up,  I suspect you'll enjoy this one.

Best part, don't need to market *any of it* _as *NFT*_.

The utility wins are the marketing, when the vehicle evolves beyond the .jpg.

The NFT is NFT; the tech just works, in background.
Actual Sentiment: positive
Predicted Sentiment polarity: 0.625
------------------------------------------------------------
REVIEW: Contributing to this canvas was one of the most enjoyable experiences I've had on #SolanaNFTs. Cheers to @sainteclectic and team for bringing u
Actual Sentiment: positive
Predicted Sentiment polarity: 0.23333333333333336
------------------------------------------------------------
REVIEW: @relay_chain $RELAY
@RocketVault_ $RVF
@RoyaleFNBR $ROYA
@ShadowsNet $DOWS
@ShieldFinance $SHLD
@solana $SOL
@sovifinance $SOVI
@splytcore $SHOPX
@SpoolFi $SPOOL
@SuperFarmDAO $SUPER
@SushiSwap $SUSHI
@standardweb3 $STND
@StarterXyz $START
@terra_money $LUNA $UST
Actual Sentiment: positive
Predicted Sentiment polarity: 0.3333333333333333
------------------------------------------------------------
REVIEW: @briahsdad @solana @UpholdInc @coinbase You could use some help

Here👇 write the support team, they 'll assist you

https://t.co/0TX5Jfprjm
Actual Sentiment: positive
Predicted Sentiment polarity: 0.0
------------------------------------------------------------
REVIEW: @SnapFingersDAO @joinrepublic @solana @RatioFinance @ProjectGalaxyHQ good project
```

```
Actual Sentiment: positive
Predicted Sentiment polarity: 0.7
-------------------------------------------------------------
```

Predict sentiment for test dataset

```python
sentiment_polarity = [textblob.TextBlob(review).sentiment.polarity for review in test_reviews]


# adapted codo to get positive score above 0, instead of score >= 0.1
predicted_sentiments = ['positive' if score > 0 else 'negative' for score in sentiment_polarity]
```

Evaluate model performance

```python
labels = ['negative', 'positive']
print(classification_report(test_sentiments, predicted_sentiments))
pd.DataFrame(confusion_matrix(test_sentiments, predicted_sentiments), index=labels, columns=labels)
```

```
              precision    recall  f1-score   support

    negative       0.20      0.54      0.29        13
    positive       0.97      0.87      0.92       222

    accuracy                           0.86       235
   macro avg       0.58      0.71      0.61       235
weighted avg       0.93      0.86      0.88       235
```

	negative	positive
negative	7	6
positive	28	194

Sentiment Analysis with AFINN

The AFINN lexicon is perhaps one of the simplest and most popular lexicons and can be used extensively for sentiment analysis. Developed and curated by Finn Årup Nielsen, you can find more details on this lexicon in the paper by Finn Årup Nielsen, entitled "A New ANEW: Evaluation of a Word List for Sentiment Analysis in Microblogs," from the proceedings of the ESWC2011 workshop.

The current version of the lexicon is AFINN-en-165. txt and it contains over 3,300 words with a polarity score associated with each word. You can find this lexicon at the author's official GitHub repository along with previous versions of this lexicon including AFINN-111 at https://github.com/ fnielsen/afinn/blob/master/afinn/data/.

The author has also created a nice wrapper library on top of this in Python called afinn, which we will be using for our analysis needs.

```
from afinn import Afinn

afn = Afinn(emoticons=True)
```

▾ Predict sentiment for sample reviews

```
for review, sentiment in zip(test_reviews[sample_review_ids], test_sentiments[sample_review_ids]):
    print('REVIEW:', review)
    print('Actual Sentiment:', sentiment)
    print('Predicted Sentiment polarity:', afn.score(review))
    print('-'*60)
```

```
 REVIEW: @Brewtoshi @solana @brave @anniekatrina @BrendanEich ^^^^ scroll on up,  I suspect you'll enjoy this one.

 Best part, don't need to market *any of it* _as *NFT*_.

 The utility wins are the marketing, when the vehicle evolves beyond the .jpg.

 The NFT is NFT; the tech just works, in background.
Actual Sentiment: positive
Predicted Sentiment polarity: 10.0
------------------------------------------------------------
 REVIEW: Contributing to this canvas was one of the most enjoyable experiences I've had on #SolanaNFTs. Cheers to @sainteclectic and team for bringing u
Actual Sentiment: positive
Predicted Sentiment polarity: 8.0
------------------------------------------------------------
 REVIEW: @relay_chain $RELAY
@RocketVault_ $RVF
@RoyaleFNBR $ROYA
@ShadowsNet $DOWS
@ShieldFinance $SHLD
@solana $SOL
@sovifinance $SOVI
@splytcore $SHOPX
@SpoolFi $SPOOL
@SuperFarmDAO $SUPER
@SushiSwap $SUSHI
@standardweb3 $STND
@StarterXyz $START
@terra_money $LUNA $UST
Actual Sentiment: positive
Predicted Sentiment polarity: 3.0
------------------------------------------------------------
```

```
REVIEW: @briahsdad @solana @UpholdInc @coinbase You could use some help

Here👇 write the support team, they 'll assist you

https://t.co/0TX5Jfprjm
Actual Sentiment: positive
Predicted Sentiment polarity: 4.0
-----------------------------------------------------------
REVIEW: @SnapFingersDAO @joinrepublic @solana @RatioFinance @ProjectGalaxyHQ good project
Actual Sentiment: positive
Predicted Sentiment polarity: 3.0
-----------------------------------------------------------
```

▾ Predict sentiment for test dataset

```
sentiment_polarity = [afn.score(review) for review in test_reviews]
# adapted codo to get positive score above 0, instead of score >= 1.0
predicted_sentiments = ['positive' if score > 0 else 'negative' for score in sentiment_polarity]
```

▾ Evaluate model performance

```
labels = ['negative', 'positive']
print(classification_report(test_sentiments, predicted_sentiments))
pd.DataFrame(confusion_matrix(test_sentiments, predicted_sentiments), index=labels, columns=labels)
```

	precision	recall	f1-score	support
negative	0.33	0.69	0.45	13
positive	0.98	0.92	0.95	222
accuracy			0.91	235
macro avg	0.66	0.81	0.70	235
weighted avg	0.94	0.91	0.92	235

	negative	positive
negative	9	4
positive	18	204

```
import nltk
import re
import numpy as np
import contractions
```

```
import contractions

stop_words = nltk.corpus.stopwords.words('english')

def normalize_document(doc):
    # lower case and remove special characters\whitespaces
    doc = re.sub(r'[^a-zA-Z0-9\s]', '', doc, re.I|re.A)
    doc = doc.strip()
    doc = contractions.fix(doc)
    return doc

normalize_corpus = np.vectorize(normalize_document)

norm_corpus = normalize_corpus(test_reviews)
len(norm_corpus)
```

```
    235
```

Sentiment Analysis with VADER

The VADER lexicon, developed by C.J. Hutto, is based on a rule-based sentiment analysis framework, specifically tuned to analyze sentiments in social media. VADER stands for Valence Aware Dictionary and sEntiment Reasoner. Details about this framework can be read in the original paper by Hutto, C.J., and Gilbert, E.E. (2014), entitled "VADER: A Parsimonious Rule-based Model for Sentiment Analysis of Social Media Text," from the proceedings of the Eighth International Conference on Weblogs and Social Media (ICWSM-14). You can use the library based on NLTK's interface under the `nltk. sentiment.vader` module.

You can also download the actual lexicon or install the framework from https:// github.com/cjhutto/vaderSentiment, which also contains detailed information about VADER. This lexicon, present in the file titled vader_lexicon.txt, contains necessary sentiment scores associated with words, emoticons, and slangs (like wtf, lol, nah, and so on).

There were a total of over 9,000 lexical features from which over 7,500 curated lexical features were finally selected in the lexicon with proper validated valence scores.

Each feature was rated on a scale from `"[-4] Extremely Negative"` to `"[4] Extremely Positive"`, with allowance for `"[0] Neutral (or Neither, N/A)"`.

The process of selecting lexical features was done by keeping all features that had a non-zero mean rating and whose standard deviation was less than 2.5, which was determined by the aggregate of ten independent raters. We depict a sample from the VADER lexicon as follows:

```
 :( -1.9 1.13578 [-2, -3, -2, 0, -1, -1, -2, -3, -1, -4]
 :) 2.0 1.18322 [2, 2, 1, 1, 1, 1, 4, 3, 4, 1]
 ...
```

```
terrorizing -3.0 1.0 [-3, -1, -4, -4, -4, -3, -2, -3, -2, -4]
thankful 2.7 0.78102 [4, 2, 2, 3, 2, 4, 3, 3, 2, 2]
```

Each line in the preceding lexicon sample depicts a unique term, which can either be an emoticon or a word. The first token indicates the word/emoticon, the second token indicates the mean sentiment polarity score, the third token indicates the standard deviation, and the final token indicates a list of scores given by 10 independent scorers.

```python
from nltk.sentiment.vader import SentimentIntensityAnalyzer

def analyze_sentiment_vader_lexicon(review,
                                    threshold=0.1,
                                    verbose=False):
    # analyze the sentiment for review
    analyzer = SentimentIntensityAnalyzer()
    scores = analyzer.polarity_scores(review)
    # get aggregate scores and final sentiment
    agg_score = scores['compound']
    final_sentiment = 'positive' if agg_score >= threshold\
                                  else 'negative'
    if verbose:
        # display detailed sentiment statistics
        positive = str(round(scores['pos'], 2)*100)+'%'
        final = round(agg_score, 2)
        negative = str(round(scores['neg'], 2)*100)+'%'
        neutral = str(round(scores['neu'], 2)*100)+'%'
        sentiment_frame = pd.DataFrame([[final_sentiment, final, positive,
                                         negative, neutral]],
                                       columns=pd.MultiIndex(levels=[['SENTIMENT STATS:'],
                                                                     ['Predicted Sentiment', 'Polarity Score',
                                                                      'Positive', 'Negative', 'Neutral']],
                                                             codes=[[0,0,0,0,0],[0,1,2,3,4]]))
        print(sentiment_frame)

    return final_sentiment
```

▾ Predict sentiment for sample reviews

```python
for review, sentiment in zip(norm_corpus[sample_review_ids], test_sentiments[sample_review_ids]):
    print('REVIEW:', review)
    print('Actual Sentiment:', sentiment)
    pred = analyze_sentiment_vader_lexicon(review, threshold=0.4, verbose=True)
    print('-'*60)
```

```
REVIEW: Brewtoshi solana brave anniekatrina BrendanEich  scroll on up  I suspect you will enjoy this one

Best part do not need to market any of it as NFT

The utility wins are the marketing when the vehicle evolves beyond the jpg

The NFT is NFT the tech just works in background
Actual Sentiment: positive
    SENTIMENT STATS:
  Predicted Sentiment Polarity Score Positive Negative Neutral
0          positive          0.92    24.0%     4.0%   72.0%
-------------------------------------------------------------
REVIEW: Contributing to this canvas was one of the most enjoyable experiences I Have had on SolanaNFTs Cheers to sainteclectic and team for bringing us
Actual Sentiment: positive
    SENTIMENT STATS:
  Predicted Sentiment Polarity Score Positive Negative Neutral
0          positive          0.86    18.0%     0.0%   82.0%
-------------------------------------------------------------
REVIEW: relaychain RELAY
RocketVault RVF
RoyaleFNBR ROYA
ShadowsNet DOWS
ShieldFinance SHLD
solana SOL
sovifinance SOVI
splytcore SHOPX
SpoolFi SPOOL
SuperFarmDAO SUPER
SushiSwap SUSHI
standardweb3 STND
StarterXyz START
terramoney LUNA UST
Actual Sentiment: positive
    SENTIMENT STATS:
  Predicted Sentiment Polarity Score              Positive Negative Neutral
0          positive          0.68 14.000000000000002%     0.0%   86.0%
-------------------------------------------------------------
REVIEW: briahsdad solana UpholdInc coinbase You could use some help

Here write the support team they ll assist you

httpstco0TX5Jfprjm
Actual Sentiment: positive
    SENTIMENT STATS:
  Predicted Sentiment Polarity Score Positive Negative Neutral
0          positive          0.66    24.0%     0.0%   76.0%
-------------------------------------------------------------
REVIEW: SnapFingersDAO joinrepublic solana RatioFinance ProjectGalaxyHQ good project
Actual Sentiment: positive
    SENTIMENT STATS:
  Predicted Sentiment Polarity Score Positive Negative Neutral
0          positive          0.44    33.0%     0.0%   67.0%
-------------------------------------------------------------
```

Predict sentiment for test dataset

```
predicted_sentiments = [analyze_sentiment_vader_lexicon(review, threshold=0.4, verbose=False) for review in test_reviews]
```

Evaluate model performance

```
labels = ['negative', 'positive']
print(classification_report(test_sentiments, predicted_sentiments))
pd.DataFrame(confusion_matrix(test_sentiments, predicted_sentiments), index=labels, columns=labels)
```

```
              precision    recall  f1-score   support

    negative       0.12      0.46      0.20        13
    positive       0.96      0.81      0.88       222

    accuracy                           0.79       235
   macro avg       0.54      0.64      0.54       235
weighted avg       0.92      0.79      0.84       235
```

	negative	positive
negative	6	7
positive	42	180

✓ 0s   completed at 12:35 PM   ● ✕